

Deep Learning

18 Variational Autoencoder

Dr. Konda Reddy Mopuri
Dept. of AI, IIT Hyderabad
Jan-May 2023

- ① Designed to reproduce input, especially reproduce the input from a learned encoding

- ① Designed to reproduce input, especially reproduce the input from a learned encoding
- ② We attempted to project the data into the latent space and model it via a probability distribution


- ① Designed to reproduce input, especially reproduce the input from a learned encoding
- ② We attempted to project the data into the latent space and model it via a probability distribution
- ③ This wasn't satisfying

Variational Autoencoders

- ① 'Regularized' autoencoder to enforce latent space 'organization'


Variational Autoencoders

- ① Key idea is to make both Encoder and Decoder stochastic
 - instead of encoding an i/p as a single point, we encode it as a distribution over the latent space




Variational Autoencoders

- ① Key idea is to make both Encoder and Decoder stochastic
 - instead of encoding an i/p as a single point, we encode it as a distribution over the latent space
- ② Latent variable z is drawn from a probability distribution for the given input x



Variational Autoencoders

- 1 Then, the reconstruction is chosen probabilistically from the sampled z




- ① Takes i/p and returns the parameters of a probability density (e.g. Gaussian, mean and covariance matrix)

- ① Takes i/p and returns the parameters of a probability density (e.g. Gaussian, mean and covariance matrix)
- ② We can sample this to get random values of the latent variable z

- ① Takes i/p and returns the parameters of a probability density (e.g. Gaussian, mean and covariance matrix)
- ② We can sample this to get random values of the latent variable z
- ③ NN implementation of the encoder gives (for every input x) a vector mean and a diagonal covariance

VAE Encoder




- 1 Decoder takes the latent vector z and returns the parameters for a distribution


- ① Decoder takes the latent vector z and returns the parameters for a distribution
- ② $p_{\phi}(x/z)$ gives mean and variance for each pixel in the output

- ① Decoder takes the latent vector z and returns the parameters for a distribution
- ② $p_{\phi}(x/z)$ gives mean and variance for each pixel in the output
- ③ Reconstruction of x is via sampling (with some assumptions, the data sample can be output)

VAE Decoder



VAE Forward pass



VAE loss function

- ① Loss for AE: l_2 distance between the input and its reconstruction

VAE loss function

- ① Loss for AE: l_2 distance between the input and its reconstruction
- ② In case of VAE: we need to learn parameters of two probability distributions

VAE loss function



- ① Loss for AE: l_2 distance between the input and its reconstruction
- ② In case of VAE: we need to learn parameters of two probability distributions
- ③ For each input x_i we maximize expected value of returning x_i (or, minimize the NLL)

$$-\mathbb{E}_{z \sim q_\theta(z/x_i)}[\log p_\phi(x_i/z)]$$

$$-\mathbb{E}_{z \sim q_{\theta}(z/x_i)}[\log p_{\phi}(x_i/z)]$$

- ① Problem: Input images may be memorized in the latent space
- similar inputs may get different representations in z space
 - close points in the latent space should not give two completely different contents once decoded

VAE loss function



$$-\mathbb{E}_{z \sim q_{\theta}(z/x_i)}[\log p_{\phi}(x_i/z)]$$

- ① Continuity and Completeness: We prefer continuous latent representations to give meaningful parameterization (e.g. smooth transition between i/ps)

$$-\mathbb{E}_{z \sim q_\theta(z/x_i)} [\log p_\phi(x_i/z)]$$

- ① Continuity and Completeness: We prefer continuous latent representations to give meaningful parameterization (e.g. smooth transition between i/ps)
- ② Solution: Force $q_\theta(z/x_i)$ to be close to a standard distribution (e.g. Gaussian)

VAE loss function


$$l_i(\theta, \phi) = -\mathbb{E}_{z \sim q_\theta(z/x_i)}[\log p_\phi(x_i/z)] + \mathbb{KL}(q_\theta(z/x_i) || p(z))$$

- ① First term promotes recovery, second term keeps encoding continuous (beats memorization)

VAE loss function

$$l_i(\theta, \phi) = -\mathbb{E}_{z \sim q_\theta(z/x_i)}[\log p_\phi(x_i/z)] + \mathbb{KL}(q_\theta(z/x_i) || p(z))$$


① Problem: Differentiating over θ and ϕ



VAE loss function

$$l_i(\theta, \phi) = -\mathbb{E}_{z \sim q_\theta(z/x_i)} [\log p_\phi(x_i/z)] + \text{KL}(q_\theta(z/x_i) || p(z))$$

- ① Reparameterization: Draw samples from $N(0,1) \rightarrow$ doesn't depend on parameters



Generation with VAE

- Sample z from the prior $p(z)$

Generation with VAE

- Sample z from the prior $p(z)$
- Run z through the decoder (ϕ) \rightarrow distribution over data

Generation with VAE

- Sample z from the prior $p(z)$
- Run z through the decoder (ϕ) \rightarrow distribution over data
- Sample from that distribution to generate the sample x

Generation with VAE



Figure credits: Wojceich

Generation with VAE




Figure credits: Kingma et al.

Edit/Manipulate samples with VAE

